

Data-Driven Methods For Engineers

Linear Regression: Task 1 & 2

Coursework 2 – Group 61



The University of Manchester
Aerospace Engineering
AERO40041 - Data-Driven Methods for Engineers - Saleh Rezaeiravesh
Semester 1, 2025

Juan Ignacio Doval Roque, Sacha Muller
09.12.2025 - final



Contents

1	Linear Regression using SGD & MSE	3
1.1	Code: LinearRegression_SGD.py	3
1.2	Console Output	6
1.3	Plots	7
2	Linear Regression using GD & MAE	8
2.1	Code: LinearRegression_MAE.py	8
2.2	Console Output	12
2.3	Plots	13
3	Contributions	14



1 | Linear Regression using SGD & MSE

1.1 Code: LinearRegression_SGD.py

The code produced for Task 1 can be found below:

```

1 # Linear Regression using Stochastic Gradient Descent (SGD)
2 # AER040041 Coursework 2 - Task 1
3 #
4 # This file was prepared by Juan Ignacio Doval Roque [10752534]
5 #
6 # NOTE: Plots have been hard coded for given data (Qdot vs dT), therefore, if other data
7 # is used,
# the plotting section may need to be adjusted accordingly. Also the path to the data
8 # file should be changed.
9
10 import numpy as np
11 import pandas as pd
12 import matplotlib.pyplot as plt
13
14 # Change as required
15 your_path = '' # Currently at root
16 data_file = "window_heat.csv"
17 y_label = '$\dot{Q}$(W)'
18 x_label = 'ΔT (°C)'
19
20 def load_and_prep_data(filename):
21     """
22         Load dataset from CSV and prepare feature matrix X and target vector y.
23         Adds a column of ones to X, for the bias term.
24     """
25     data = pd.read_csv(filename)
26
27     X = data.iloc[:, :-1].values
28     y = data.iloc[:, -1].values.reshape(-1, 1)
29
30     ones = np.ones((X.shape[0], 1))
31     X = np.concatenate((ones, X), axis=1)
32
33     return X, y

```

Listing 1 - Information, Imports, Paths, and Data preparation

```

1 def mse_loss(X, y, w):
2     """
3         Compute Mean Squared Error loss.
4
5         MSE = (1/N) * sum((y_pred - y)^2)
6     """
7     N = X.shape[0]
8     y_pred = X @ w
9     error = y_pred - y
10    loss = (1/N) * np.sum(error ** 2)
11
12    return loss

```

Listing 2 - MSE Loss function



```

1 def sgd_linear_regression(X, y, learning_rate=0.01, num_epochs=100, seed=42):
2     """
3         Train linear regression using Stochastic Gradient Descent.
4
5         In SGD, we update weights after each individual sample.
6         One EPOCH means we have processed every sample exactly once.
7
8         Parameters:
9         -----
10        X : numpy array of shape (N, D+1)
11            Feature matrix with bias column
12        y : numpy array of shape (N, 1)
13            Target values
14        learning_rate : float
15            Step size for gradient updates
16        num_epochs : int
17            Number of complete passes through the dataset
18        seed : int
19            Random seed for reproducibility
20
21     Returns:
22     -----
23        w : numpy array
24            Learned weights
25        loss_history : list
26            Loss value recorded at the end of each epoch
27        """
28    np.random.seed(seed)
29
30    N, D = X.shape
31    w = np.random.randn(D, 1) * 0.01
32    loss_history = []
33
34    for epoch in range(num_epochs):
35        indices = np.random.permutation(N)
36
37        for i in indices:
38            xi = X[i:i+1, :]
39            yi = y[i:i+1, :]
40            y_pred = xi @ w
41            error = y_pred - yi
42            gradient = 2 * xi.T @ error
43            w = w - learning_rate * gradient
44
45            current_loss = mse_loss(X, y, w)
46            loss_history.append(current_loss)
47
48            if (epoch + 1) % 10 == 0 or epoch == 0:
49                print(f"Epoch {epoch + 1}/{num_epochs}, Loss: {current_loss:.6f}")
50
51    return w, loss_history
52
53 def predict(X, w):
54     """Make predictions using learned weights."""
55     return X @ w

```

Listing 3 - Stochastic Gradient Descent Linear Regression function and Prediction function



```

1  if __name__ == "__main__":
2
3      X, y = load_and_prep_data(your_path + data_file)
4
5      dT_original = X[:, 1].copy()
6
7      print("DATA CHECK:")
8      print(f"dT range: {dT_original.min():.2f} to {dT_original.max():.2f} °C")
9      print(f"{y_label} range: {y.min():.2f} to {y.max():.2f} W")
10     print(f"Dataset shape: {X.shape[0]} samples, {X.shape[1]-1} features")
11
12    X_normalized = X.copy()
13    means = X[:, 1:].mean(axis=0)
14    stds = X[:, 1:].std(axis=0)
15    stds[stds == 0] = 1
16    X_normalized[:, 1:] = (X[:, 1:] - means) / stds
17
18    print("\nStart of SGD Linear Regression:")
19
20    w_sgd, loss_history = sgd_linear_regression(
21        X_normalized, y,
22        learning_rate=0.01,
23        num_epochs=100,
24        seed=42
25    )
26
27    print(f"\nFinal Loss (MSE): {loss_history[-1]:.6f}")
28    print(f"Learned weights (normalized): {w_sgd.flatten()}")
29
30    y_pred = predict(X_normalized, w_sgd)
31
32    final_mse = np.mean((y - y_pred) ** 2)
33    final_rmse = np.sqrt(final_mse)
34    print(f"Final RMSE: {final_rmse:.6f}")

```

Listing 4 - Main Execution Block

```

1  fig, axes = plt.subplots(3, 1, figsize=(8, 12))
2
3  # Plot 1: Loss vs Epoch
4  axes[0].plot(range(1, len(loss_history) + 1), loss_history, 'b-', linewidth=2)
5  axes[0].set_xlabel('Epoch', fontsize=12)
6  axes[0].set_ylabel('MSE Loss', fontsize=12)
7  axes[0].set_title('SGD Training: Loss vs Epoch', fontsize=14)
8  axes[0].grid(True, alpha=0.3)
9
10 # Plot 2: Qdot vs dT with fitted line
11 axes[1].scatter(dT_original, y, alpha=0.6, edgecolors='black',
12                  linewidth=0.5, label='Training data points', color='blue')
13 dT_line = np.linspace(dT_original.min(), dT_original.max(), 100)
14 dT_line_normalized = (dT_line - means[0]) / stds[0]
15 X_line = np.column_stack([np.ones(100), dT_line_normalized])
16 y_line = X_line @ w_sgd
17
18 axes[1].plot(dT_line, y_line, 'r-', linewidth=2, label='Fitted line')
19 axes[1].set_xlabel(x_label, fontsize=12)
20 axes[1].set_ylabel(y_label, fontsize=12)
21 axes[1].set_title('Heat Loss vs Temperature Difference', fontsize=14)
22 axes[1].legend()
23 axes[1].grid(True, alpha=0.3)
24 #Indented

```

Listing 5 - Plotting pt.1



```

1      # Plot 3: Predictions vs Actual
2      axes[2].scatter(y, y_pred, alpha=0.6, edgecolors='black',
3                          linewidth=0.5, label='Training data points', color='blue')
4      min_val = min(y.min(), y_pred.min())
5      max_val = max(y.max(), y_pred.max())
6      axes[2].plot([min_val, max_val], [min_val, max_val], 'r--',
7                      linewidth=2, label='Prediction')
8      axes[2].set_xlabel(f'Actual {y_label}', fontsize=12)
9      axes[2].set_ylabel(f'Predicted {y_label}', fontsize=12)
10     axes[2].set_title('Predictions vs Actual Values', fontsize=14)
11     axes[2].legend()
12     axes[2].grid(True, alpha=0.3)
13
14     plt.tight_layout()
15     plt.savefig(your_path + 'LinearRegression_SGD_output.png', dpi=150,
16                  bbox_inches='tight')
16     plt.show()
17 #Indented

```

Listing 6 - Plotting pt.2

1.2 Console Output

```

1 DATA CHECK:
2 dT range: 1.05 to 21.93 °C
3 $\dot{Q}$(W) range: 482.65 to 11507.79 W
4 Dataset shape: 24 samples, 1 features
5
6 Start of SGD Linear Regression:
7 Epoch 1/100, Loss: 19754458.888643
8 Epoch 10/100, Loss: 595389.836357
9 Epoch 20/100, Loss: 594480.745186
10 Epoch 30/100, Loss: 594339.488119
11 Epoch 40/100, Loss: 594665.175322
12 Epoch 50/100, Loss: 596306.570341
13 Epoch 60/100, Loss: 595375.779382
14 Epoch 70/100, Loss: 595024.150064
15 Epoch 80/100, Loss: 594531.145023
16 Epoch 90/100, Loss: 594669.893471
17 Epoch 100/100, Loss: 594472.436477
18
19 Final Loss (MSE): 594472.436477
20 Learned weights (normalized): [6402.42813965 3115.93548705]
21 Final RMSE: 771.020387
22
23 Plot saved as 'LinearRegression_SGD_output.png'

```

Listing 7 - Console Output: Data Ranges, Dataset Shape, Epochs, Loss, Weights, Plot path



1.3 Plots

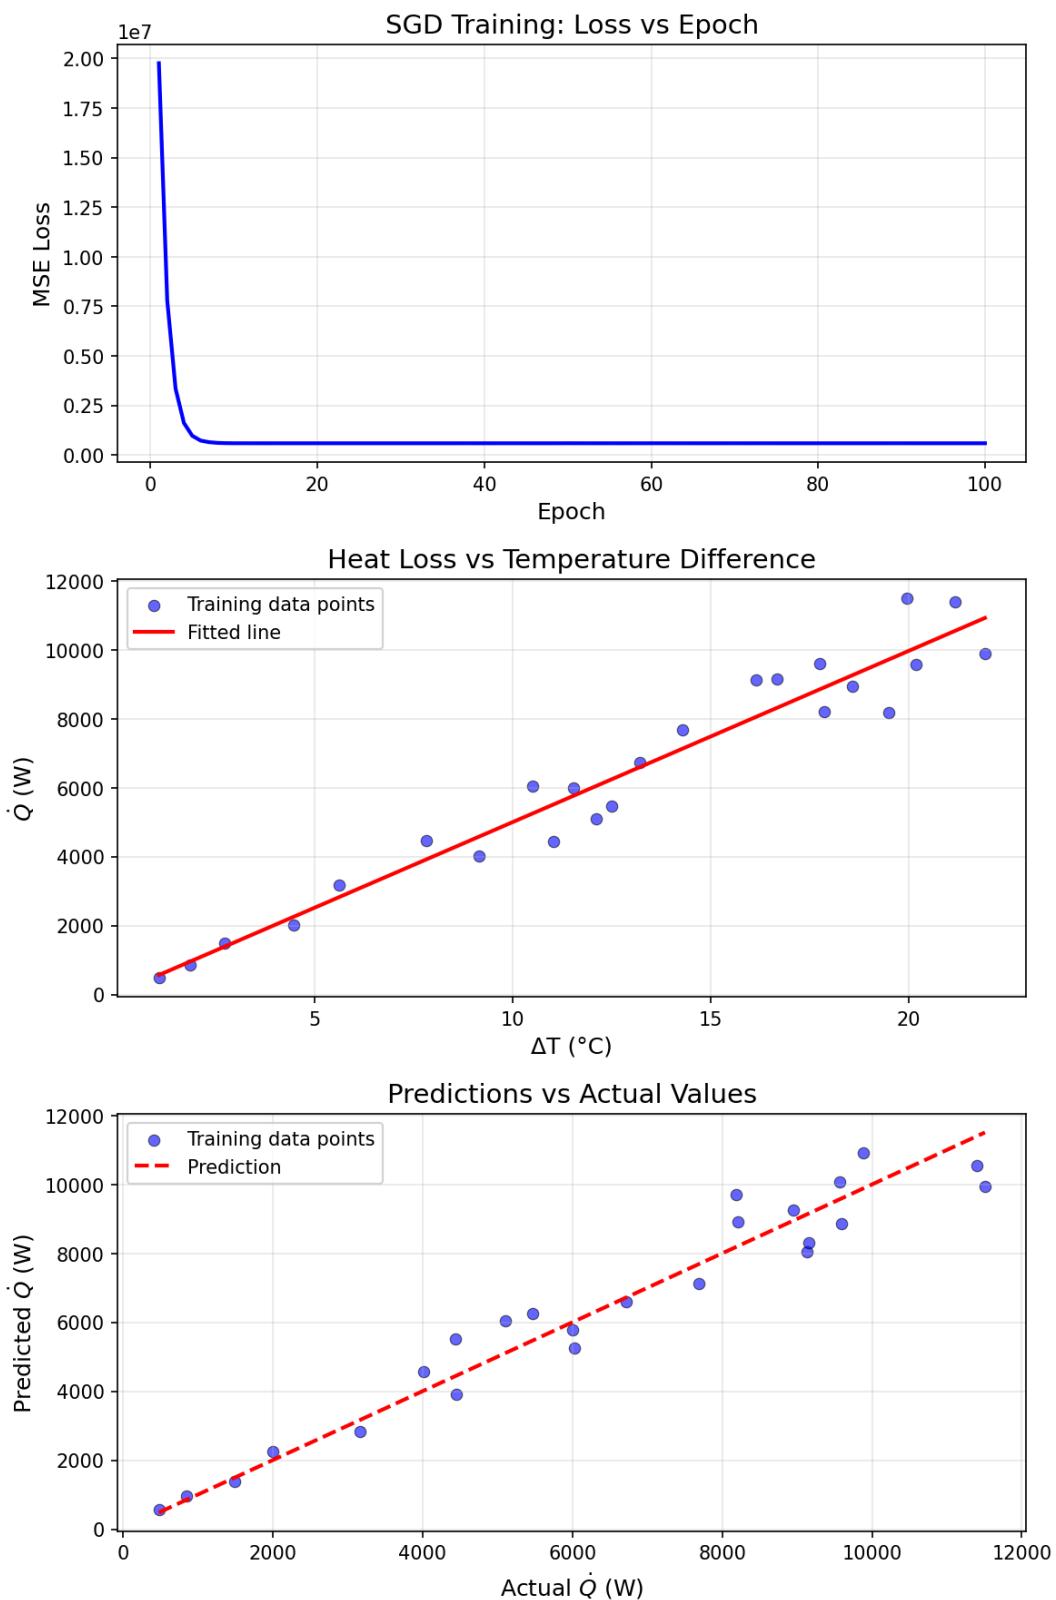


Figure 1 - SGD with MSE: (Top) Loss vs Epoch, (Middle) Heat Loss vs Temperature Difference, (Bottom) Prediction of Target (W) vs Actual Values from data file



2 | Linear Regression using GD & MAE

2.1 Code: LinearRegression_MAE.py

The code produced for Task 2 can be found below:

```

1 # Linear Regression using Mean Absolute Error (MAE) Loss
2 # AER040041 Coursework 2 - Task 2
3 #
4 # This file was prepared by Juan Ignacio Doval Roque [10752534]
5 #
6 # NOTE: Plots have been hard coded for given data (Qdot vs dT), therefore, if other data
7 is used,
# the plotting section may need to be adjusted accordingly. Also the path to the data
8 file should be changed.
9
10 import numpy as np
11 import pandas as pd
12 import matplotlib.pyplot as plt
13
14 your_path = '' # Currently at root
15 data_file = "window_heat.csv"
16 y_label = '$\dot{Q}$(W)'
17 x_label = 'ΔT (°C)'
18
19 def load_and_prep_data(filename):
20     """
21         Load dataset from CSV and prepare feature matrix X and target vector y.
22         Adds a column of ones to X for the bias term.
23     """
24     data = pd.read_csv(filename)
25
26     X = data.iloc[:, :-1].values
27     y = data.iloc[:, -1].values.reshape(-1, 1)
28
29     ones = np.ones((X.shape[0], 1))
30     X = np.concatenate((ones, X), axis=1)
31
32     return X, y

```

Listing 8 - Information, Imports, Paths, and Data preparation

```

1 def mae_loss(X, y, w):
2     """
3         Compute Mean Absolute Error loss.
4
5         MAE = (1/N) * sum(|y_pred - y|)
6
7         MAE is more robust to outliers than MSE because it doesn't square errors.
8     """
9     N = X.shape[0]
10    y_pred = X @ w
11    error = y_pred - y
12    loss = (1/N) * np.sum(np.abs(error))
13    return loss

```

Listing 9 - MAE Loss function



```

1  def mae_gradient(X, y, w):
2      """
3          Compute the gradient of MAE loss with respect to weights.
4
5          The gradient of |x| is sign(x), where:
6              sign(x) = +1 if x > 0
7              sign(x) = -1 if x < 0
8              sign(x) = 0 if x = 0
9
10         Therefore:
11             d(MAE)/dw_j = (1/N) * sum(sign(y_pred - y) * x_j)
12
13         In matrix form:
14             gradient = (1/N) * X^T @ sign(y_pred - y)
15             """
16
17         N = X.shape[0]
18         y_pred = X @ w
19         error = y_pred - y
20
21         gradient = (1/N) * X.T @ np.sign(error)
22
23     return gradient

```

Listing 10 - MSE gradient computation



```

1 def gd_mae(X, y, learning_rate=0.01, num_iterations=1000):
2     """
3         Train linear regression using Gradient Descent with MAE loss.
4
5         This is BATCH gradient descent: we use all samples for each update.
6
7         Parameters:
8         -----
9         X : numpy array of shape (N, D+1)
10            Feature matrix with bias column
11         y : numpy array of shape (N, 1)
12            Target values
13         learning_rate : float
14            Step size for gradient updates
15         num_iterations : int
16            Number of gradient descent steps
17
18         Returns:
19         -----
20         w : numpy array
21            Learned weights
22         loss_history : list
23            MAE loss value at each iteration
24         """
25         N, D = X.shape
26
27         w = np.zeros((D, 1))
28         w[0] = np.mean(y)
29
30         loss_history = []
31         best_loss = float('inf')
32         patience = 500
33         no_improve = 0
34
35         for iteration in range(num_iterations):
36             gradient = mae_gradient(X, y, w)
37
38             w = w - learning_rate * gradient
39
40             current_loss = mae_loss(X, y, w)
41             loss_history.append(current_loss)
42
43             if current_loss < best_loss - 1e-6:
44                 best_loss = current_loss
45                 no_improve = 0
46             else:
47                 no_improve += 1
48
49             if no_improve >= patience:
50                 print(f"Early stopping at iteration {iteration + 1}")
51                 break
52
53             if (iteration + 1) % 500 == 0 or iteration == 0:
54                 print(f"Iteration {iteration + 1}/{num_iterations}, MAE Loss:
55 {current_loss:.6f}")
56
57         return w, loss_history
58
59     def predict(X, w):
60         """Make predictions using learned weights."""
61         return X @ w

```

Listing 11 - Gradient Descent Linear Regression function and Prediction function



```

1  if __name__ == "__main__":
2      X, y = load_and_prep_data(your_path + data_file)
3
4      dT_original = X[:, 1].copy()
5
6      print("DATA CHECK:")
7      print(f"dT range: {dT_original.min():.2f} to {dT_original.max():.2f} °C")
8      print(f"{y_label} range: {y.min():.2f} to {y.max():.2f} W")
9      print(f"Dataset shape: {X.shape[0]} samples, {X.shape[1]-1} features")
10
11     X_normalized = X.copy()
12     means = X[:, 1:1].mean(axis=0)
13     stds = X[:, 1:1].std(axis=0)
14     stds[stds == 0] = 1
15     X_normalized[:, 1:] = (X[:, 1:] - means) / stds
16
17     print("\nTraining w/ Gradient Descent (MAE Loss):")
18
19     w_mae, loss_history_mae = gd_mae(
20         X_normalized, y,
21         learning_rate=50.0,
22         num_iterations=5000
23     )
24
25     print(f"\nFinal MAE Loss: {loss_history_mae[-1]:.6f}")
26     print(f"Learned weights (normalized): {w_mae.flatten()}")
27
28     y_pred_mae = predict(X_normalized, w_mae)
29
30     final_mae = np.mean(np.abs(y - y_pred_mae))
31     final_mse = np.mean((y - y_pred_mae) ** 2)
32     final_rmse = np.sqrt(final_mse)
33     print(f"\nFinal MAE: {final_mae:.6f}")
34     print(f"Final MSE: {final_mse:.6f}")
35     print(f"Final RMSE: {final_rmse:.6f}")

```

Listing 12 - Main Execution Block

```

1  fig, axes = plt.subplots(3, 1, figsize=(8, 12))
2
3  # Plot 1: Loss vs Iteration
4  axes[0].plot(range(1, len(loss_history_mae) + 1), loss_history_mae, 'b-', linewidth=2)
5  axes[0].set_xlabel('Iteration', fontsize=12)
6  axes[0].set_ylabel('MAE Loss', fontsize=12)
7  axes[0].set_title('Gradient Descent with MAE: Loss vs Iteration', fontsize=14)
8  axes[0].grid(True, alpha=0.3)
9
10 # Plot 2: Qdot vs dT with fitted line
11 axes[1].scatter(dT_original, y, alpha=0.6, edgecolors='black',
12                  linewidth=0.5, label='Data points', color='blue')
13 dT_line = np.linspace(dT_original.min(), dT_original.max(), 100)
14 dT_line_normalized = (dT_line - means[0]) / stds[0]
15 X_line = np.column_stack([np.ones(100), dT_line_normalized])
16 y_line = X_line @ w_mae
17
18 axes[1].plot(dT_line, y_line, 'r-', linewidth=2, label='Fitted line (MAE)')
19 axes[1].set_xlabel(x_label, fontsize=12)
20 axes[1].set_ylabel(y_label, fontsize=12)
21 axes[1].set_title('Heat Loss vs Temperature Difference', fontsize=14)
22 axes[1].legend()
23 axes[1].grid(True, alpha=0.3)
24 #Indented

```

Listing 13 - Plotting pt.1



```

1      # Plot 3: Predictions vs Actual
2      axes[2].scatter(y, y_pred_mae, alpha=0.6, edgecolors='black',
3                          linewidth=0.5, label='Training data points', color='blue')
4      min_val = min(y.min(), y_pred_mae.min())
5      max_val = max(y.max(), y_pred_mae.max())
6      axes[2].plot([min_val, max_val], [min_val, max_val], 'r--',
7                      linewidth=2, label='Prediction')
8      axes[2].set_xlabel(f'Actual {y_label}', fontsize=12)
9      axes[2].set_ylabel(f'Predicted {y_label}', fontsize=12)
10     axes[2].set_title('MAE Model: Predictions vs Actual Values', fontsize=14)
11     axes[2].legend()
12     axes[2].grid(True, alpha=0.3)
13
14     plt.tight_layout()
15     plt.savefig(your_path + 'LinearRegression_MAE_output.png', dpi=150,
16                 bbox_inches='tight')
17     plt.show()
18
19     print(f"\nPlot saved as '{your_path}LinearRegression_MAE_output.png'")
20 #Indented

```

Listing 14 - Plotting pt.2

2.2 Console Output

```

1  DATA CHECK:
2  dT range: 1.05 to 21.93 °C
3  $\dot{Q}$ (W) range: 482.65 to 11507.79 W
4  Dataset shape: 24 samples, 1 features
5
6  Training w/ Gradient Descent (MAE Loss):
7  Iteration 1/5000, MAE Loss: 2734.424588
8  Iteration 500/5000, MAE Loss: 645.291521
9  Early stopping at iteration 941
10
11 Final MAE Loss: 645.346541
12 Learned weights (normalized): [6523.29413845 3217.08664574]
13
14 Final MAE: 645.346541
15 Final MSE: 617986.882721
16 Final RMSE: 786.121417
17
18 Plot saved as 'LinearRegression_MAE_output.png'

```

Listing 15 - Console Output: Data Ranges, Dataset Shape, Epochs, Loss, Weights, Plot path



2.3 Plots

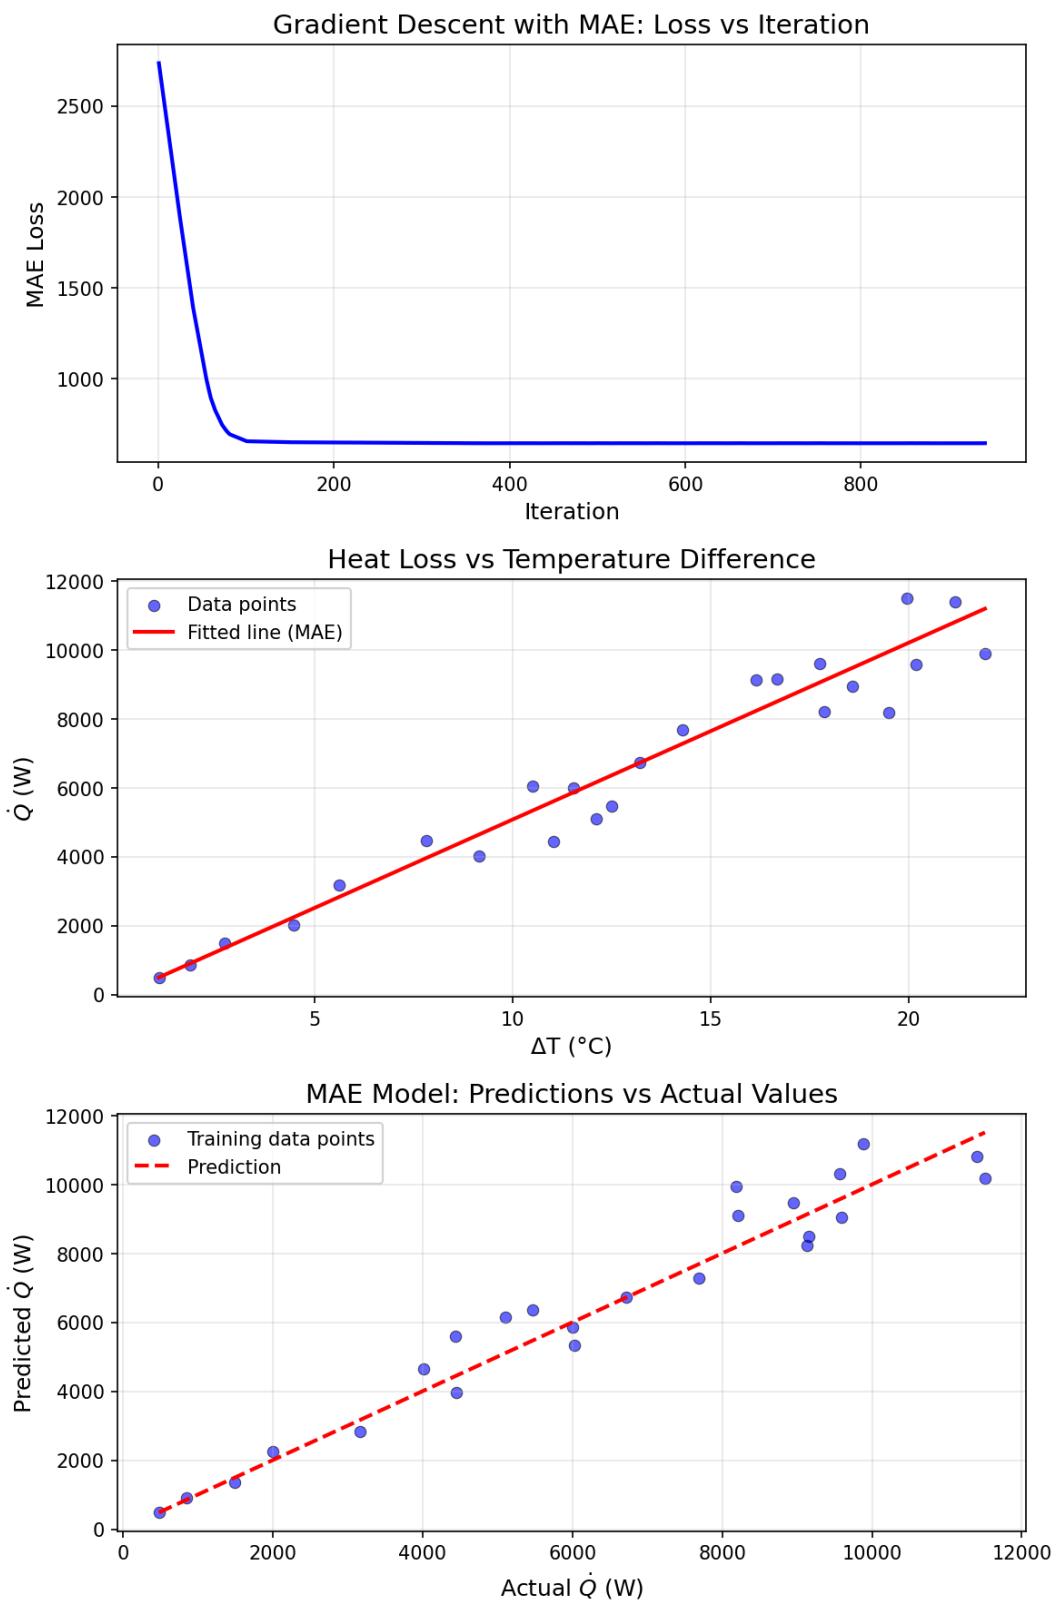


Figure 2 - GD with MAE: (Top) Loss vs Epoch, (Middle) Heat Loss vs Temperature Difference, (Bottom) Prediction of Target (W) vs Actual Values from data file



3 | Contributions

Juan created and ran the code for Tasks 1 & 2; Linear Regression. This pdf was created by Juan.

Sacha created and ran the code for Task 3; Neural Network for classification. Sacha provides his pdf separately.